

XenGT Setup Guide

October 2013

Contents

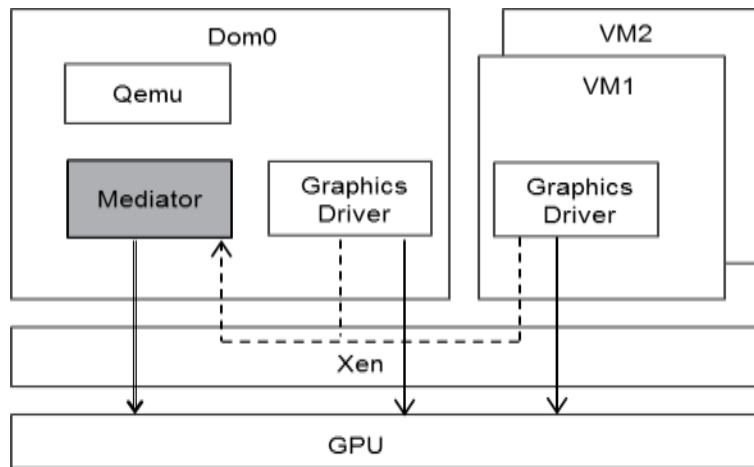
1	Introduction	4
2	System Requirements	5
2.1	Operating System Requirements	5
2.2	Hardware Requirements	5
2.3	Software Requirements	5
2.3.1	Proxy Configuration Setup	5
2.3.2	Install Basic Packages in Ubuntu	5
3	Build and Install Instructions	7
3.1	Source Repositories	7
3.2	Building Kernel.....	7
3.3	Building Xen and Qemu	7
3.4	Grub Setup.....	8
3.5	Dom0 System Config Setup for XenGT	9
3.5.1	Starting Xen Services by Default	9
3.5.2	Building Graphic Stack – Optional (Not needed for Ubuntu 13.04)	9
3.5.3	Configuring Xen Bridge.....	11
3.6	Guest Setup.....	11
3.6.1	General Setup	11
3.6.2	Guest Config File	11
3.7	Linux Guest System Setup for XenGT.....	12
3.7.1	Kernel and Modules Update.....	12
3.7.2	Disabling Grub Graphics Mode.....	13
3.7.3	Graphics Stack Update for Linux Guest – Optional (Not needed for Ubuntu 13.04)	14
4	VM Life Cycle Management	15
4.1	Guest Creation	15
4.2	Listing Information about Domains	15

4.3	Guest Destroy	15
5	XenGT Control Interfaces	16
5.1	Display Switch	16
5.2	Xen API.....	16
6	Features Supported	17
6.1	Per-VM Features.....	17
6.2	Virtualization Features	17
7	Known Issues	18
8	Appendix	19
8.1	Licenses	19
8.1.1	General Statement	19
8.1.2	Licenses of the Newly-created Files.....	19

1 Introduction

The Graphics Processing Unit (GPU) has become a fundamental building block in today's computing environment, accelerating tasks ranging from entertaining (gaming, video playback, etc.), GUI acceleration, office applications (such as CAD, photoshop) and high performance computing. Recently, there observes a trend toward adding GPU accelerations to virtual machines provided by popular desktop virtualization. In the meantime, there are also demands for buying GPU computing resources from the cloud. GPU virtualization is becoming a demanding and challenging industry asking.

Intel has brought its answer to this challenge with **XenGT**, which is a mediated pass-through solution based on Intel Gen Graphics hardware using the well-known Xen hypervisor. As illustrated below, XenGT allows running a native graphics driver in VMs to achieve high



————> Pass-through - - - -> Trap ———> Hypercall

performance. Each VM is allowed to access a partial performance critical resource without hypervisor intervention. Privileged operations are trapped and forwarded to the mediator for emulation. The mediator creates a virtual GPU context for each VM and schedules one of them to run on physical Gen graphics hardware. In our implementation, the mediator is a separate driver residing in Dom0 kernel, called **vgt driver**.

Currently, XenGT supports 4 accelerated VMs (Dom0 + 3 HVM DomU) running together.

We've verified XenGT's functionality using the 64-bit version of Ubuntu 12.04 and 13.04.

2 *System Requirements*

2.1 Operating System Requirements

The build and install environment has been validated in using x86_64 Ubuntu 12.04 and 13.04 as host.

2.2 Hardware Requirements

4th generation Intel Processor Graphics is required.

2.3 Software Requirements

2.3.1 Proxy Configuration Setup

If you building package behind firewall, you would need to setup following proxies in order to download needed libraries

```
# export http_proxy=<proxy_server>:<proxy_port>
# export https_proxy=<proxy_server>:<proxy_port>
# export ftp_proxy=<proxy_server>:<proxy_port>
```

To configure the Proxy for apt, you need add following lines into /etc/apt/apt.conf

```
Acquire::http::proxy "<proxy_server>:<proxy_port>";
Acquire::https::proxy "<proxy_server>:<proxy_port>";
Acquire::ftp::proxy "<proxy_server>:<proxy_port>";
```

2.3.2 Install Basic Packages in Ubuntu

```
# apt-get update
# apt-get install libarchive-dev libghc-bzlib-dev libghc6-bzlib-dev \
zlib1g-dev mercurial gettext bcc iasl uuid-dev libncurses5-dev kpartx \
libperl-dev libgtk2.0-dev libc6-dev-i386 libaio-dev libstd1.2-dev \
nfs-common libyajl-dev libx11-dev autoconf libtool xsltproc bison flex \
```

```
xutils-dev xserver-xorg-dev x11proto-gl-dev libx11-xcb-dev vncviewer \  
libxcb-glx0 libxcb-glx0-dev libxcb-dri2-0-dev libxcb-xfixes0-dev bridge-utils \  
python-dev bin86 git vim libssl-dev pciutils-dev tightvncserver ssh texinfo -y
```

3 *Build and Install Instructions*

3.1 Source Repositories

Xen: <https://github.com/01org/XenGT-Preview-xen>

Linux: <https://github.com/01org/XenGT-Preview-kernel.git>

Qemu: <https://github.com/01org/XenGT-Preview-qemu.git>

3.2 Building Kernel

```
# git clone https://github.com/01org/XenGT-Preview-kernel.git
# cd XenGT-Preview-kernel/
# git clone git://kernel.ubuntu.com/ubuntu/linux.git linux-vgt
# cd linux-vgt
# git checkout -b v3.8.13.4 v3.8.13.4
# patch -p1 < ../linux-vgt.patch
# cp config-3.8-dom0 .config
# make -j8 && make modules_install
# mkinitramfs -o /boot/initrd-vgt-3.8.13.4-vgt.img 3.8.13.4-vgt+
# cp arch/x86/boot/bzImage /boot/vmlinuz-vgt-3.8.13.4-vgt
# cp vgt.rules /etc/udev/rules.d
```

3.3 Building Xen and Qemu

```
# git clone https://github.com/01org/XenGT-Preview-xen.git
# cd XenGT-Preview-xen/
# git clone https://github.com/01org/XenGT-Preview-qemu.git
# git clone git://xenbits.xen.org/xen.git xen-vgt
# git clone git://git.qemu-project.org/qemu.git qemu-xen
# cd qemu-xen
# git checkout -b v1.3.0 v1.3.0
# patch -p1 < ../XenGT-Preview-qemu/qemu-vgt.patch
# cd ../xen-vgt
```

```
# git checkout -b RELEASE-4.3 RELEASE-4.3
# patch -p1 < ../XenGT-Preview-xen/xen-vgt.patch
# cp -r ../qemu-xen tools/
# sed -i
'/QEMU_UPSTREAM_URL/s:http\:\/\/xenbits.xen.org/git-http/qemu-upstream-4.3-testin
g.git:${(XEN_ROOT)}/tools/qemu-xen:' Config.mk
# sed -i
'/QEMU_UPSTREAM_URL/s:git\:\/\/xenbits.xen.org/qemu-upstream-4.3-testing.git:${(XEN_ROOT)}/tools/qemu-xen:' Config.mk
# ./autogen.sh
# ./configure --prefix=/usr # XEN4.3 changes the default path to /usr/local
# make -j8 xen tools
# cp xen/xen.gz /boot/xen-vgt.gz
# make install-tools PYTHON_PREFIX_ARG=
# rm /etc/ld.so.conf.d/lib64.conf
# ldconfig
```

3.4 Grub Setup

You need manually add a new grub entry in /boot/grub/grub.cfg and make the entry as the default one when booting. Below is a reference grub entry for you.

```
UUID (2e01a442-d848-4695-b031-9296ce3105b1) and root partition (hd0, msdos1)
below are just reference which should be updated according to the user's environment.
menuentry 'Xen-VGT 3.8.13.4' --class ubuntu --class gnu-linux --class gnu --class os {
    insmod part_msdos
    insmod ext2
    set root='(hd0,msdos1)'
    search --no-floppy --fs-uuid --set=root 2e01a442-d848-4695-b031-9296ce3105b1
    multiboot /boot/xen-vgt.gz dom0_mem=2048M loglvl=all guest_loglvl=all
    conring_size=4M noreboot
    module /boot/vmlinuz-vgt-3.8.13.4-vgt
    root=UUID=2e01a442-d848-4695-b031-9296ce3105b1 rw rd_NO_LUKS rd_NO_LVM
    LANG=en_US.UTF-8 rd_NO_MD SYSFONT=latarcyrheb-sun16 rhgb crashkernel=auto
```



```
KEYBOARDTYPE=pc KEYTABLE=us rd_NO_DM ignore_loglevel console=tty0
console=hvc0 consoleblank=0 log_buf_len=4M xen_vgt.hvm_boot_foreground=1
module /boot/initrd-vgt-3.8.13.4-vgt.img
```

Description of suggested parameters for grub

Configuration option	Description
xen_vgt.hvm_boot_foreground	Make VM immediately visible on the screen, after creation
xen_vgt.vgt	Option to enable/disable vgt for Dom0. 0 to disable vgt, 1 to enable vgt(default value)

3.5 Dom0 System Config Setup for XenGT

3.5.1 Starting Xen Services by Default

```
# update-rc.d xencommons defaults
```

3.5.2 Building Graphic Stack – Optional (Not needed for Ubuntu 13.04)

If you want to run 3D workloads in Dom0, the user mode driver update is required to support 4th generation Intel Processor Graphics. It is not a required step if dom0 does not start X.

```
# apt-get install git build-essential libtool autoconf libpthread-stubs0-dev \
libpciaccess-dev xutils-dev xserver-xorg-dev bison x11proto-gl-dev \
x11proto-xext-dev libxdamage-dev xserver-xorg-dev libx11-xcb-dev \
libxcb-glx0-dev libxcb-dri2-0-dev libxext-dev libexpat1-dev libxcb-xfixes0-dev
```

If you want to build and install the libraries to system "/usr/" directory directly, you do not need below environment variables. Just simply add option "--prefix=/usr" for "autogen.sh" command. If you do not want to pollute system, please follow below steps:

```
# export LD_LIBRARY_PATH=/opt/hsw/usr/lib
# export
PKG_CONFIG_PATH=/opt/hsw/usr/lib/pkgconfig:/opt/hsw/usr/share/pkgconfig/
# mkdir -p /opt/hsw/usr
# cd /opt/hsw
```

```
# git clone git://anongit.freedesktop.org/git/mesa/drm
# cd drm
# git checkout 171666e4b8127c17c68ea0d44cf4e81ec342f2d0
# ./autogen.sh --prefix=/opt/hsw/usr
# make && make install
# cd ..
# git clone
git://anongit.freedesktop.org/git/xorg/driver/xf86-video-intel
# cd xf86-video-intel
# git checkout b6d2bb961517
# ./autogen.sh --prefix=/opt/hsw/usr
# make && make install
# cd ..
```

Notice that in above steps we check out a specific revision of x driver. The reason not to use the latest commit is that the driver has dependence to X. Latest x driver requires you to build latest Xorg as well. Above commit has proved to work fine with default Ubuntu 12.04 Xorg.

```
# apt-get build-dep mesa
# git clone git://anongit.freedesktop.org/git/mesa/mesa
# cd mesa
# git checkout a585b8f3a6681d1138ed1a33ed4c3195a53c2a73
# ./autogen.sh --disable-gallium-egl --disable-gallium-gbm --without-gallium-drivers
--with-dri-drivers=i965 --prefix=/opt/hsw/usr
# make && make install
# cd ..
```

Then use the new driver for your system:

```
# cd /usr/lib/xorg/modules/drivers
backup original intel_drv.so
# ln -sf /opt/hsw/usr/lib/xorg/modules/drivers/intel_drv.so intel_drv.so
# cp newGL.conf /etc/ld.so.conf.d/
newGL.conf contains only one line of "/opt/hsw/usr/lib"
# ldconfig
```

3.5.3 Configuring Xen Bridge

After Dom0 reboots, run following commands to make a bridge "xenbr0" for guest network. (Assume the IP address of system could be acquired via DHCP and the default network interface with network connection is "eth0")

```
# brctl addbr xenbr0
# ifconfig eth0 0.0.0.0 down
# brctl addif xenbr0 eth0
# ifconfig eth0 up
# dhclient xenbr0
```

3.6 Guest Setup

3.6.1 General Setup

You need create an empty image with at least 10GB for guest. Here we take Ubuntu 13.04 guest image as example.

```
# dd if=/dev/zero of=system-10G.img bs=1M seek=10000 count=0
```

After you get the installation ISO for Ubuntu 13.04 guest, you could set the xmexample.conf by following section 3.6.2 with 3 changes and start the guest with the config file.

```
disk = [ 'file:/path/system-10G.img,hda,w', 'file:/path/to/ubuntu/iso,hdc:cdrom,r' ]
boot="d"
vgt=0
```

Then you could follow the normal procedure to finish the Ubuntu installation.

3.6.2 Guest Config File

You could copy the xmexample.conf from /etc/xen and modify the parameters as following.

```
Kernel = "hvmloader"
builder = 'hvm'
memory = 2048
name = "vgtHVMDomain"
vif = [ 'type=ioemu, bridge=xenbr0' ]
disk = [ 'file:/path/to/image/file,hda,w', ',hdc:cdrom,r' ]
```

```

#device_model = 'qemu-dm'
device_model_version='qemu-xen'
device_model_override='/usr/lib/xen/bin/qemu-system-i386'
sdl=1
opengl=1
vnc=0
vncpasswd=""
serial='pty'
tsc_mode=0
stdvga = 0
usb=1
usbdevice='tablet'
keymap='en-us'
vgt=1
vgt_low_gm_sz=64
vgt_high_gm_sz=448
vgt_fence_sz=4

```

Description of parameters specific for XenGT

Configuration option	Description
vgt	Enable virtual graphics
vgt_low_gm_sz	The low gm size which is CPU visible. For linux guest, it should be at least 64MB
vgt_high_gm_sz	The high gm size which is CPU invisible
vgt_fence_sz	The number of the fence registers, default is 4

3.7 Linux Guest System Setup for XenGT

3.7.1 Kernel and Modules Update

Assume that you have an Ubuntu x86_64 13.04 image. You could update the kernel and user mode drivers in guest image with following commands:

```
# kpartx -a -v ./ubuntu.img
```

The output will be something like below:

```
add map loop0p1 (253:0): 0 29638656 linear /dev/loop0 2048
add map loop0p2 (253:1): 0 1075202 linear /dev/loop0 29642750
add map loop0p5 : 0 1075200 linear 253:1 2
```

Mount loop0p1 to /mnt:

```
# mount /dev/mapper/loop0p1 /mnt/
```

Follow the steps in Section 3.2 to build Dom0 kernel and modules in guest with "chroot".

```
# chroot /mnt/
# exit
```

Alternatively, you could copy the kernel/initrd and modules from Dom0 to guest directly:

```
# cp /boot/vmlinuz-vgt-3.8.13.4-vgt /mnt/boot/
# cp /boot/initrd-vgt-3.8.13.4-vgt.img /mnt/boot/
# cp -r /lib/modules/3.8.13.4-vgt+ /mnt/lib/modules
```

Then you could run commands as below to fresh the guest image:

```
# umount /mnt
# kpartx -d -v ./ubuntu.img
```

3.7.2 Disabling Grub Graphics Mode

Due to a known issue for XenGT (see section 7), grub graphics mode should be disabled.

You could boot up the guest with "vgt=0" in xmexample.conf and update the "/etc/default/grub" with following change:

```
GRUB_TERMINAL=console
```

Then run

```
# update-grub
```

Then you should add one new entry in "/boot/grub/grub.cfg" for the new kernel and initrd added in section 3.7.1. Now, the image is ready for XenGT. You need set "vgt=1" in the xmexample.conf and start the guest.

3.7.3 Graphics Stack Update for Linux Guest – Optional (Not needed for Ubuntu 13.04)

If you want to run 3D workloads in Linux guest, the user mode driver update is required to support 4th generation Intel Processor Graphics. You could follow the steps in section 3.5.2 to build the Graphics Stack for Linux guest.

4 *VM Life Cycle Management*

4.1 Guest Creation

To create a Guest, you need a configure file mentioned in section 3.6.2 and use the command:

```
# xl create xmexample.hvm
```

4.2 Listing Information about Domains

To get the information (ID, name, vCPU, Mem, etc) of all domains, you could use “xl list”.

```
# xl list
```

Name	ID	Mem	VCPU	State	Time(s)
Domain-0	0	730	8	r-----	1139.2
ExampleHVMDomain	7	2048	1	r-----	1.0

4.3 Guest Destroy

Shutting down a Guest should be triggered in Guest and follow the normal procedure for Linux. To destroy a Guest, you could use “xl destroy” with Domain ID or Name, as below:

```
# xl destroy ExampleHVMDomain
```

5 *XenGT Control Interfaces*

5.1 Display Switch

When a Guest is created successfully with vgt enabled, the monitor display will be switched to the Guest with `xen_vgt.hvm_boot_foreground` set. A sys interface is provided to switch display between domains:

```
# xl list
```

Name	ID	Mem	VCPUs	State	Time(s)
Domain-0	0	730	8	r-----	3263.5
ExampleHVMDomain	11	2048	1	r-----	4.7

```
# cat /sys/kernel/vgt/control/foreground_vm
```

```
11
```

You could change the display by echo Domain ID into the sys interface:

```
# echo 0 > /sys/kernel/vgt/control/foreground_vm
```

```
# cat /sys/kernel/vgt/control/foreground_vm
```

```
0
```

5.2 Xen API

Refer to `XenGT-API.txt` under `xen-vgt` on how to use XenGT API interfaces.

6 *Features Supported*

6.1 Per-VM Features

Features or Areas	Status
SMP Dom0 and Guest	Supported
2D Blitter	Supported
3D Rendering (Direct3D/OpenGL)	Supported
HW-accelerated Video Playback	Preliminary Supported
Single Monitor (HDMI/VGA/eDP/DP)	Supported
Multiple Monitors	Supported
PPGTT	Supported
Dom0 S3	Preliminary supported

6.2 Virtualization Features

Features or Areas	Status
Up to 3 vGT Guests	3 Guests, each with 128MB aperture.
Render Context Switch	Supported
Display Switch	It works only if all VMs use the same resolution.
VM Life Cycle	Supported
XenGT Interfaces (APIs)	Refer to the API document.
Monitor Hotplug	Supported

7 *Known Issues*

- After Dom0 S3, Ubuntu VM may hang for several minutes.
- Multiple displays hotplug might have display issue w/ VM.
- eDP panel could not display when Dom0 booted up with eDP, HDMI and DP plugged together.
- Monitor Hotplug does not work for VM on ULT platform.
- Foreground VM switch may not work sometimes.
- Resolution change is not supported.
- 2GB memory is suggested for VM to run most 3D workloads.
- Dom0 might be unstable with workload running in guest for more than 12 hours.
- Graphics mode of grub2 not working in guest:
Due to a QEMU bug, grub of Ubuntu guest cannot work on graphic mode. The solution is to disable gfxmode by having "GRUB_TERMINAL=console" in /etc/default/grub.
- compiz not working:
compiz may not work in Dom0 and linux guest. You need select "Unity 2D" during UX login instead. Or you could just edit "/etc/lightdm/lightdm.conf", change "user-session=ubuntu" to be "user-session=ubuntu-2d".
- keymap might be incorrect in guest
config file may need to explicitly specify "keymap='en-us'". Although it looks like the default value, earlier we saw the problem of wrong keymap code if it is not explicitly set.

8 *Appendix*

8.1 Licenses

8.1.1 General Statement

Only newly-created files are explicitly specified with a license here. Any changes to existing files of Xen/Linux/Qemu are subject to the licenses of the files.

8.1.2 Licenses of the Newly-created Files

8.1.2.1 Xen Side

All the newly-created files are under GPLv2:

tools/firmware/hvmloader/vgt_devtable.h

xen/arch/x86/hvm/vgt.c

xen/arch/x86/vgt.c

xen/include/asm-x86/vgt.h

8.1.2.2 Linux Side

These newly-created files are under GPLv2:

arch/x86/include/asm/xen/x86_emulate.h

arch/x86/xen/vgt_emulate.c

arch/x86/xen/x86_emulate.c

drivers/xen/vgt/debugfs.c

drivers/xen/vgt/dev.c

drivers/xen/vgt/hypercall.c

drivers/xen/vgt/klog.c

drivers/xen/vgt/Makefile

drivers/xen/vgt/sysfs.c

drivers/xen/vgt/trace.h

tools/vgt/klog.c

tools/vgt/Makefile

```
tools/vgt/README
tools/vgt/vgt_report
vgt.rules
XenGT-API.txt
```

These newly-created files are under dual GPLv2/MIT:

```
drivers/xen/vgt/aperture_gm.c
drivers/xen/vgt/cfg_space.c
drivers/xen/vgt/cmd_parser.c
drivers/xen/vgt/devtable.h
drivers/xen/vgt/display.c
drivers/xen/vgt/edid.c
drivers/xen/vgt/edid.h
drivers/xen/vgt/gtt.c
drivers/xen/vgt/handlers.c
drivers/xen/vgt/instance.c
drivers/xen/vgt/interrupt.c
drivers/xen/vgt/mmio.c
drivers/xen/vgt/reg.h
drivers/xen/vgt/render.c
drivers/xen/vgt/sched.c
drivers/xen/vgt/utility.c
drivers/xen/vgt/vgt.c
drivers/xen/vgt/vgt.h
drivers/xen/vgt/fb_decoder.c
include/xen/interface/hvm/ioreq.h
drivers/xen/vgt/cmd_parser.h
include/xen/vgt.h
include/xen/vgt-if.h
include/xen/fb_decoder.h
```

8.1.2.3 Qemu Side

There are 2 newly-created files, which are under dual GPLv2/MIT:

```
hw/vga-xengt.c
hw/vga-xengt.h
```